DASC7606
2025/11/06

Assignment II: Machine Translation

Weijian Deng (kendwj@hku.hk)

Shu Chen (schen59@hku.hk)

Find it on Github: https://github.com/hkukend/DASC7606A-B-A2

Optional Assignment:

- **This assignment is optional** due to the coming final exams
- **The final assignment grades** will be the maximum between Assignment 1 and 2

Objectives:

- **Use the Hugging Face Hub for datasets:** discover, download, and prepare translation datasets
- **Load pretrained models and tokenizers** from the community
- **Build training pipelines using** transformers, datasets, accelerate, and evaluate
- **Fine-tune a translation model** end-to-end for a chosen language pair (zh-sim->en)
- **Evaluate translation quality** using BLEU (via sacrebleu) and report results
- **Develop debugging skills** for identifying and fixing common deep learning issues in HF version

HuggingFace 🤗 :

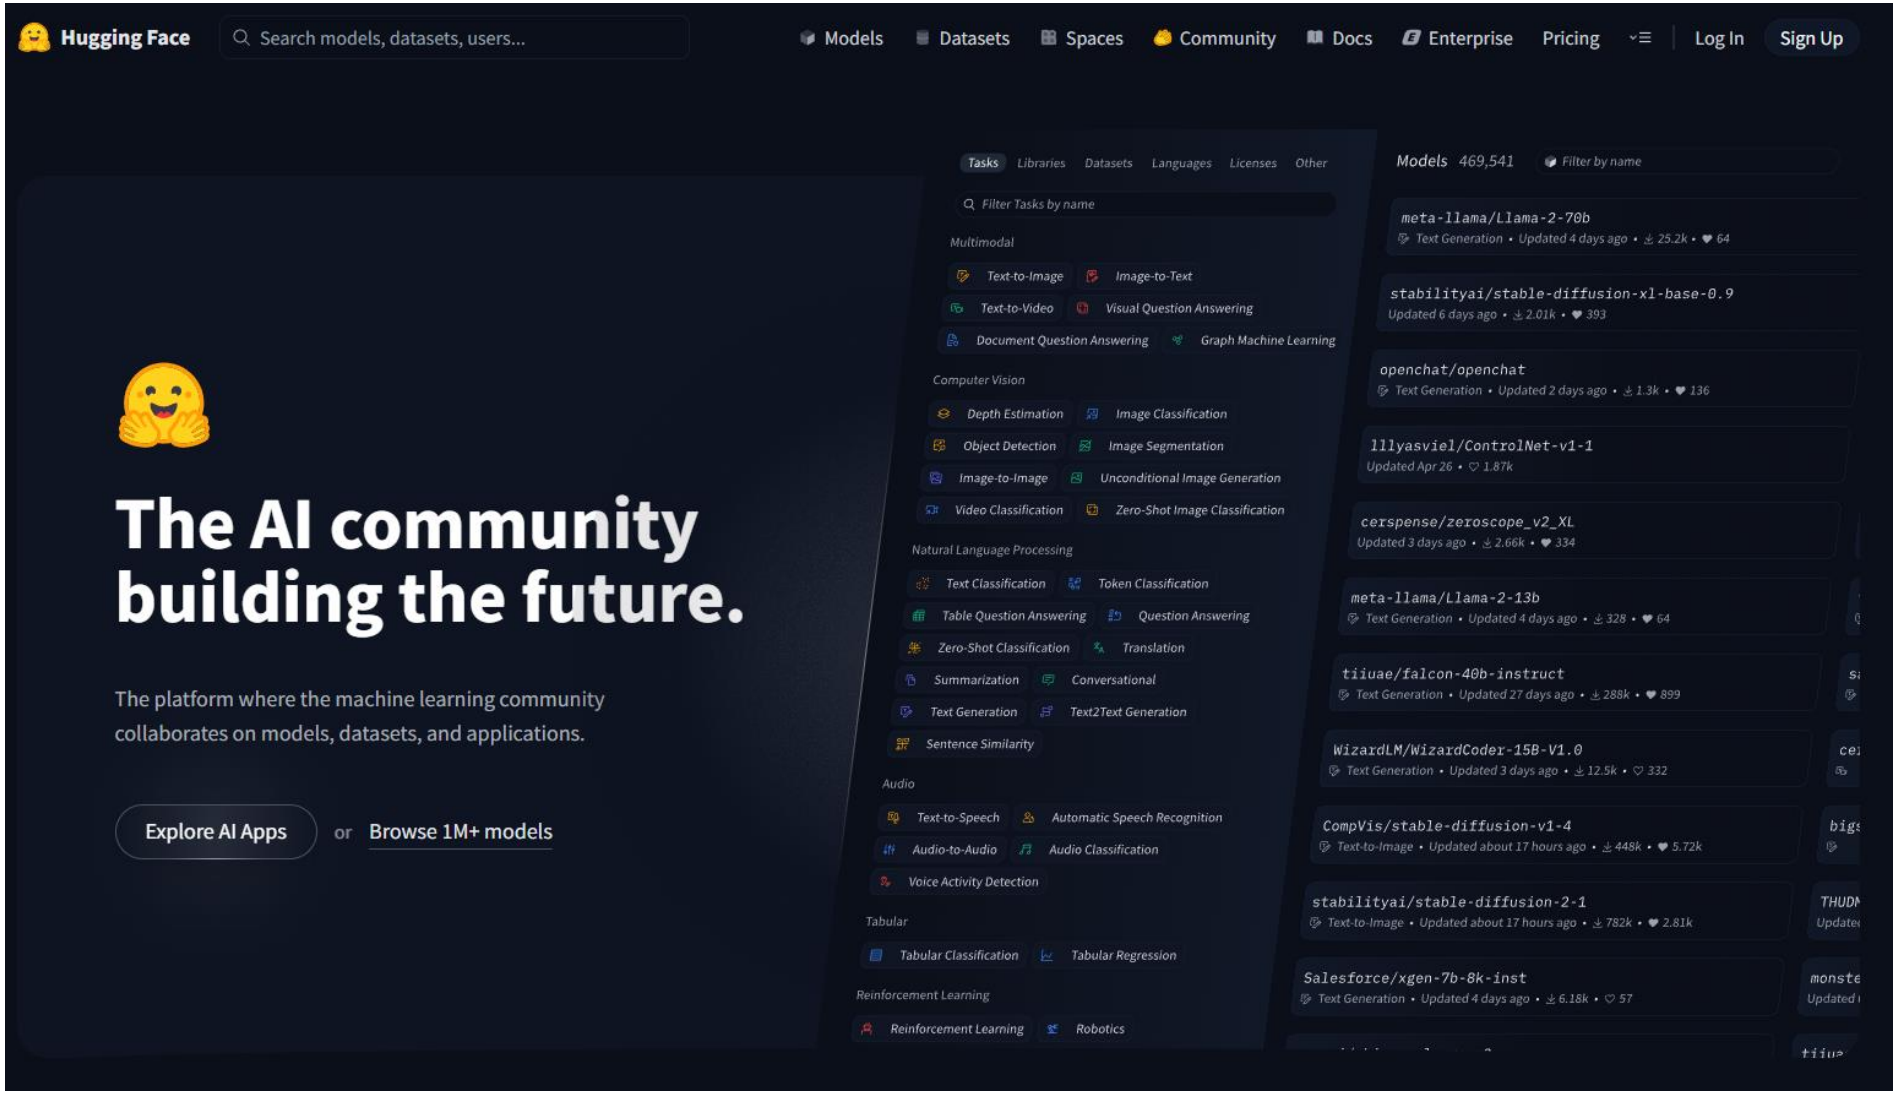

HuggingFace 🤗 Datasets:

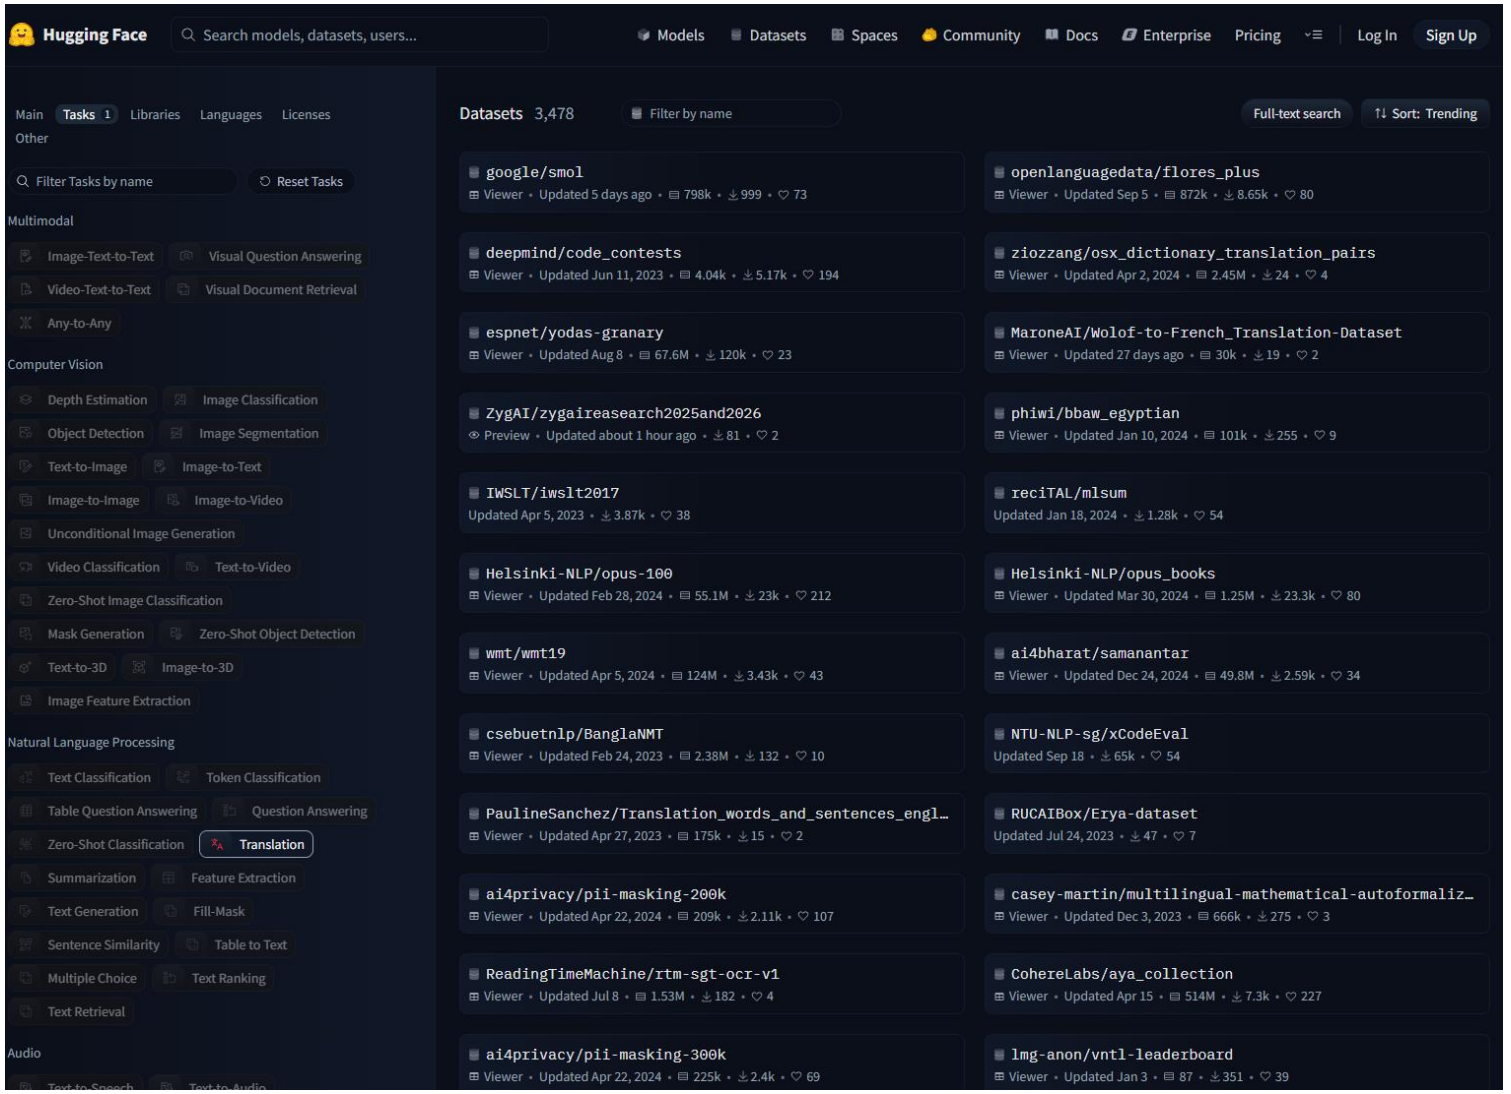

Use one line to load datasets from HugginFace platform:

```
>>> from datasets import load_dataset
>>> dataset = load_dataset("wmt19", "zh-en")
>>> dataset
DatasetDict({
    train: Dataset({
        features: ['translation'],
        num_rows: 25984574
    })
    validation: Dataset({
        features: ['translation'],
        num_rows: 3981
    })
})
```

HuggingFace 🤗 Transformers:

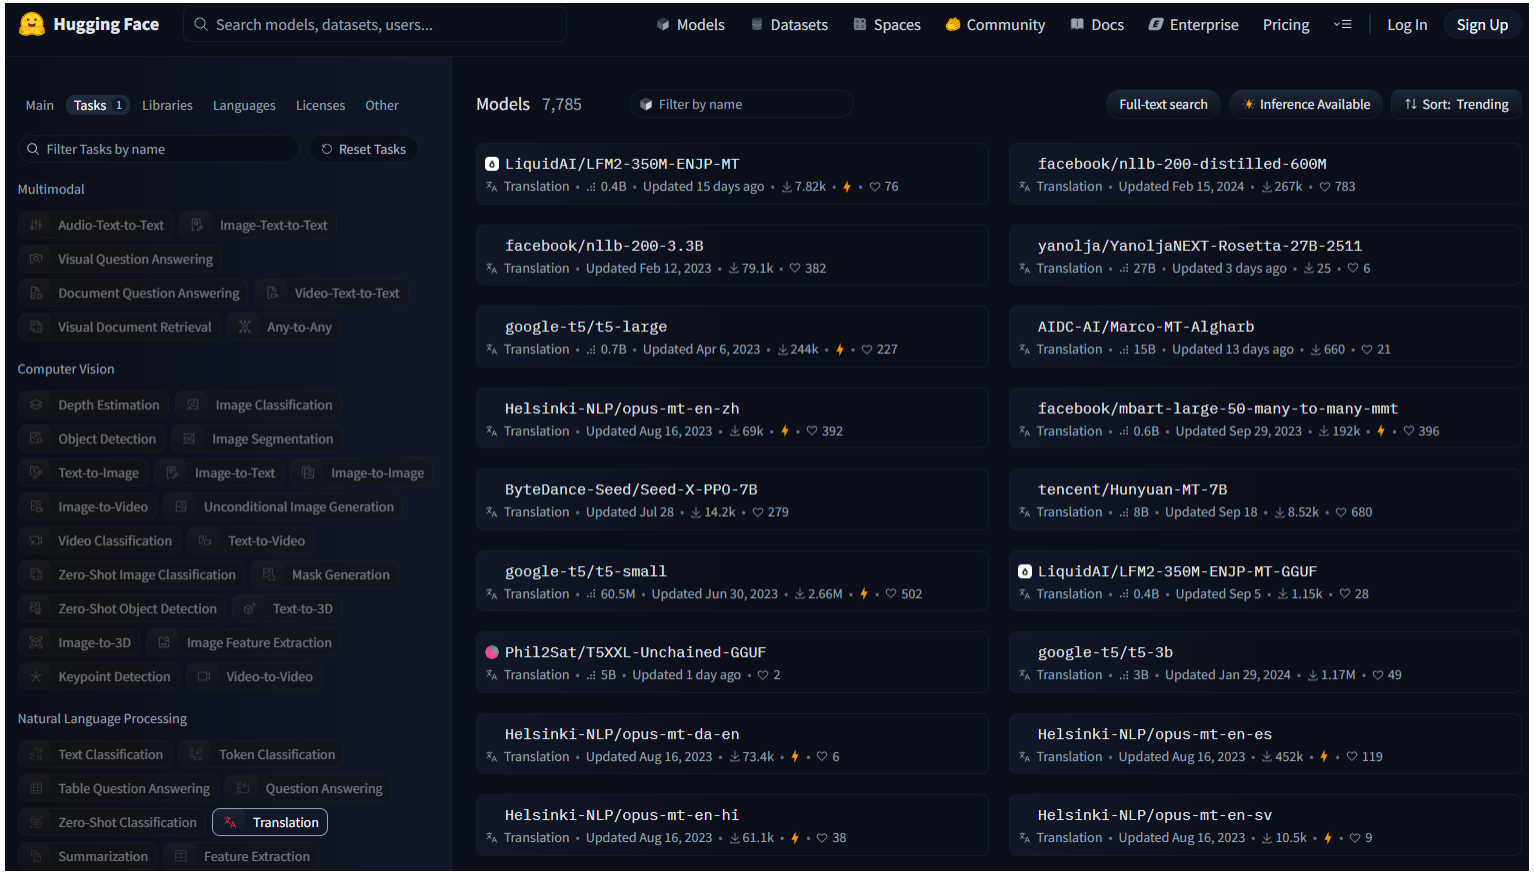

Use one line to load tokenizers and pretrained models from HugginFace platform:

```
>>> from transformers import AutoTokenizer, AutoModel
>>> tokenizer = AutoTokenizer.from_pretrained(pretrained_model_name_or_path = "Helsinki-NLP/opus-mt-en-zh")
/[redacted]/.venv/lib/python3.13/site-packages/transformers/models/marian/tokenizati
on_marian.py:175: UserWarning: Recommended: pip install sacremoses.
  warnings.warn("Recommended: pip install sacremoses.")
>>> tokenizer
MarianTokenizer(name_or_path='Helsinki-NLP/opus-mt-en-zh', vocab_size=65001, model_max_length=512, is_fast=False,
 padding_side='right', truncation_side='right', special_tokens={'eos_token': '</s>', 'unk_token': '<unk>', 'pad_t
oken': '<pad>'}, clean_up_tokenization_spaces=False, added_tokens_decoder={
        0: AddedToken("</s>", rstrip=False, lstrip=False, single_word=False, normalized=False, special=True),
        1: AddedToken("<unk>", rstrip=False, lstrip=False, single_word=False, normalized=False, special=True),
        65000: AddedToken("<pad>", rstrip=False, lstrip=False, single_word=False, normalized=False, special=True)
,
}
)
>>> model = AutoModel.from_pretrained(pretrained_model_name_or_path = "Helsinki-NLP/opus-mt-en-zh")
>>> model
MarianModel(
  (shared): Embedding(65001, 512, padding_idx=65000)
  (encoder): MarianEncoder(
    (embed_tokens): Embedding(65001, 512, padding_idx=65000)
    (embed_positions): MarianSinusoidalPositionalEmbedding(512, 512)
    (layers): ModuleList(
      (0-5): 6 x MarianEncoderLayer(
        (self_attn): MarianAttention(
          (k_proj): Linear(in_features=512, out_features=512, bias=True)
          (v_proj): Linear(in_features=512, out_features=512, bias=True)
          (q_proj): Linear(in_features=512, out_features=512, bias=True)
          (out_proj): Linear(in_features=512, out_features=512, bias=True)
        )
        (self_attn_layer_norm): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
        (activation_fn): SiLU()
        (fc1): Linear(in_features=512, out_features=2048, bias=True)
        (fc2): Linear(in_features=2048, out_features=512, bias=True)
        (final_layer_norm): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
      )
    )
  )
  (decoder): MarianDecoder(
    (embed_tokens): Embedding(65001, 512, padding_idx=65000)
    (embed_positions): MarianSinusoidalPositionalEmbedding(512, 512)
    (layers): ModuleList(
      (0-5): 6 x MarianDecoderLayer(
        (self_attn): MarianAttention(
          (k_proj): Linear(in_features=512, out_features=512, bias=True)
          (v_proj): Linear(in_features=512, out_features=512, bias=True)
          (q_proj): Linear(in_features=512, out_features=512, bias=True)
          (out_proj): Linear(in_features=512, out_features=512, bias=True)
        )
        (activation_fn): SiLU()
```

What we will do:

The workflow is organized into 5 files:

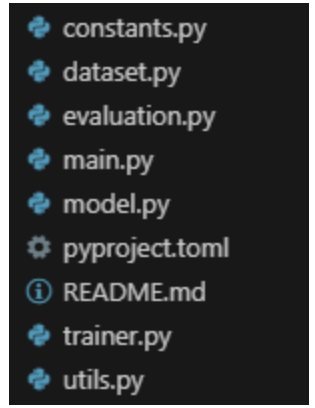

- **Dataset Sourcing & Preparation (dataset.py)**
  Explore HF Datasets, load splits, perform filtering/mapping, and train/validation/test preparation

- **Model & Tokenizer Setup (model.py)**
  Select a suitable base model from the Hub and initialize tokenizer and config

- **Training Pipeline (trainer.py)**
  Configure TrainingArguments, data collators, metrics, logging, and checkpointing

- **Run & Monitor Training (main.py)**
  Orchestrate end-to-end training and validation, with periodic evaluation

- **Evaluation & Reporting (evaluation.py)**
  Compute BLEU on the held-out test set; save artifacts and summary

What we will do:

- **Goal**: Fine-tune a suitable HF model on a translation dataset and achieve competitive BLEU.
- **What you can modify**: You may change any files in the repo **except** `main.py`, part of `evaluation.py` and `utils.py`. The test set must not be modified.
- **What to improve**:
  - **Enriching Datasets**: Use a more varied dataset (no leakage of the test set into training)
  - **Base model choice**: Select an appropriate pretrained model for your language pair
  - **Training pipeline**: Tune hyperparameters (batch size, LR, epochs, schedulers, label smoothing, gradient accumulation)
  - **Data processing**: Tokenization lengths, filtering, cleaning, language codes, special tokens
  - **Advanced HF features**: Mixed precision (`fp16`/`bf16`), gradient checkpointing, LoRA/PEFT, better data collators, scheduler choices, early stopping

Example Accepted Datasets/Models

- Datasets: `wmt14`, `wmt16`, `wmt19`, `opus100`, `ted_talks_iwslt`, etc. (via HF Datasets)
- Models: MarianMT (Helsinki-NLP/opus-mt-xx-yy), mT5, MBART-50, M2M100, NLLB-200 (ensure your pair is supported), and even **LLMs**, etc.

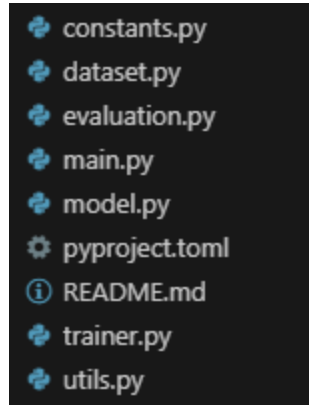

Grading:

We will re-run `main.py` and evaluate on the fixed test set. BLEU (SacreBLEU) is the primary metric.

Important Considerations:

1. **Error-Free Execution**: Your code must run without errors (and avoid GPU OOM on the provided environment)
2. **Correct Data Usage**: Do not alter or leak the test set into training
3. **No Personal Pre-trained Model**: "Downloads last month" of loaded pre-trained model on HuggingFace should be **GREATER THAN 10**
4. **Reasonable Performance**: Achieve competitive BLEU given the chosen model and setup
5. **Runtime**: Complete in a reasonable time budget (≤ 12 hours with **ONE GPU on HKU GPU Farm**)

BLEU-based Grading (dummy thresholds, subject to adjustment):

- **BLEU ≥ 25**: 100%
- **BLEU ≥ 24**: 90%
- **BLEU ≥ 23**: 80%
- **BLEU ≥ 22**: 70%
- **BLEU ≥ 21**: 60%
- **BLEU ≥ 20**: 50%
- **BLEU < 20 / Fail to reproduce / Overtime**: 0%

Important dates:

- Assignment II Release: Nov. 06 (Thursday)
- Submission Deadline: Nov. 30 (Sunday) (23:59 GMT+8)

Late submission policy :

- All submissions later than the deadline will NOT be accepted

Questions!

If any more questions, please contact kendwj@hku.hk
or schen59@hku.hk